

Semantic Chunking for RAG



Plaban Nayak · Following

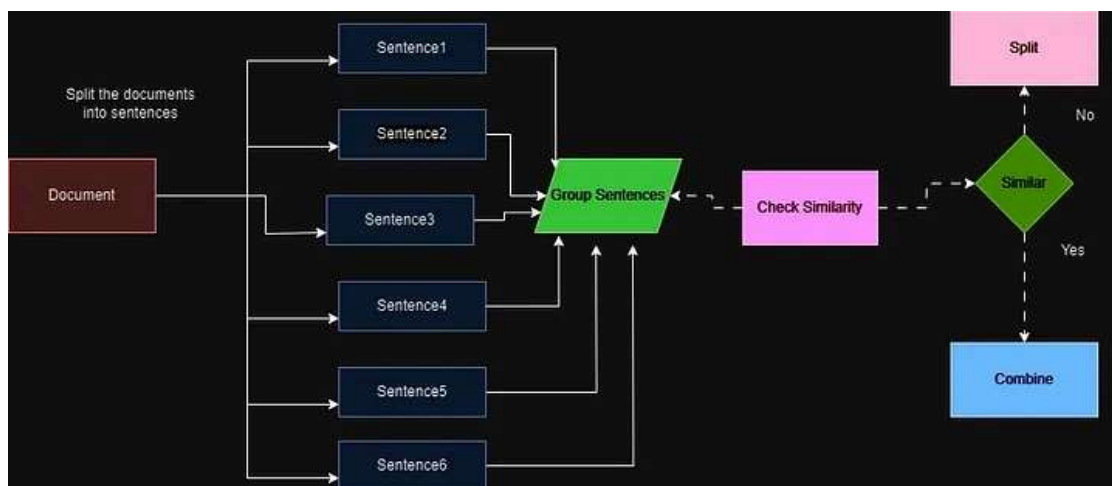
Published in The AI Forum · 17 min read · Apr 21, 2024



375



4



What is Chunking ?

In order to abide by the context window of the LLM , we usually break text into smaller parts / pieces which is called chunking.

What is RAG?

LLMs, although capable of generating text that is both meaningful and grammatically correct, these LLMs suffer from a problem called hallucination. Hallucination in LLMs is the concept where the LLMs confidently generate wrong answers, that is they make up wrong answers in a way that makes us believe that it is true. This has been a major problem since the introduction of the LLMs. These hallucinations lead to incorrect and factually wrong answers. Hence Retrieval Augmented Generation was introduced.

In RAG, we take a list of documents/chunks of documents and encode these textual documents into a numerical representation called vector embeddings, where a single vector embedding represents a single chunk of document and stores them in a database called vector store. The models required for encoding these chunks into embeddings are called encoding models or bi-encoders. These encoders are trained on a large corpus of data, thus making them powerful enough to encode the chunks of documents in a single vector embedding representation.

The retrieval greatly depends on how the chunks are manifested and stored in the vectorstore. Finding the right chunk size for any given text is a very hard question in general.

Improving Retrieval can be done by various retrieval method. But it can also be done by better chunking strategy.

Different chunking methods:

- Fixed size chunking
- Recursive Chunking
- Document Specific Chunking
- Semantic Chunking
- Agentic Chunking

Fixed Size Chunking: This is the most common and straightforward approach to chunking: we simply decide the number of tokens in our chunk and, optionally, whether there should be any overlap between them. In general, we will want to keep some overlap between chunks to make sure that the semantic context doesn't get lost between chunks. Fixed-sized chunking will be the best path in most common cases. Compared to other forms of chunking, fixed-sized chunking is computationally cheap and simple to use since it doesn't require the use of any NLP libraries.

Recursive Chunking : Recursive chunking divides the input text into smaller chunks in a hierarchical and iterative manner using a set of separators. If the

initial attempt at splitting the text doesn't produce chunks of the desired size or structure, the method recursively calls itself on the resulting chunks with a different separator or criterion until the desired chunk size or structure is achieved. This means that while the chunks aren't going to be exactly the same size, they'll still "aspire" to be of a similar size. Leverages what is good about fixed size chunk and overlap.

Document Specific Chunking: It takes into consideration the structure of the document . Instead of using a set number of characters or recursive process it creates chunks that align with the logical sections of the document like paragraphs or sub sections. By doing this it maintains the author's organization of the content thereby keeping the text coherent. It makes the retrieved information more relevant and useful, particularly for structured documents with clearly defined sections. It can handle formats such as Markdown, Html, etc.

Semantic Chunking: Semantic Chunking considers the relationships within the text. It divides the text into meaningful, semantically complete chunks. This approach ensures the information's integrity during retrieval, leading to a more accurate and contextually appropriate outcome. It is slower compared to the previous chunking strategy

Agentic Chunk: The hypothesis here is to process documents in a fashion that humans would do.

1. We start at the top of the document, treating the first part as a chunk.
2. We continue down the document, deciding if a new sentence or piece of information belongs with the first chunk or should start a new one
3. We keep this up until we reach the end of the document.

This approach is still being tested and isn't quite ready for the big leagues due to the time it takes to process multiple LLM calls and the cost of those calls. There's no implementation available in public libraries just yet.

Here we will experiment with Semantic chunking and Recursive Retriever .

Comparison of methods steps:

1. Load the Document
2. Chunk the Document using the following two methods: Semantic chunking and Recursive Retriever .
3. Assess qualitative and quantitative improvements with RAGAS

Semantic Chunks

Semantic chunking involves taking the embeddings of every sentence in the document, comparing the similarity of all sentences with each other, and then grouping sentences with the most similar embeddings together.

By focusing on the text's meaning and context, Semantic Chunking significantly enhances the quality of retrieval. It's a top-notch choice when maintaining the semantic integrity of the text is vital.

The hypothesis here is we can use embeddings of individual sentences to make more meaningful chunks. Basic idea is as follows :-

1. *Split the documents into sentences based on separators(.,?!,)*
2. *Index each sentence based on position.*
3. *Group: Choose how many sentences to be on either side. Add a buffer of sentences on either side of our selected sentence.*
4. *Calculate distance between group of sentences.*
5. *Merge groups based on similarity i.e. keep similar sentences together.*
6. *Split the sentences that are not similar.*

Technology Stack Used

- **Langchain** :LangChain is an open-source framework designed to simplify the creation of applications using large language models (LLMs). It provides a standard interface for chains, lots of integrations with other tools, and end-to-end chains for common applications.

- **LLM:** Groq's **Language Processing Unit (LPU)** is a cutting-edge technology designed to significantly enhance AI computing performance, especially for Large Language Models (LLMs). The primary goal of the Groq LPU system is to provide real-time, low-latency experiences with exceptional inference performance.
- **Embedding Model:** FastEmbed is a lightweight, fast, Python library built for embedding generation.
- **Evaluation:** Ragas offers metrics tailored for evaluating each component of your RAG pipeline in isolation.

Code Implementation

Install the required dependencies

```
langchain_community langchain ragas chromadb langchain-groq fastembed pypdf openai
```

```
langchain==0.1.16  
langchain-community==0.0.34  
langchain-core==0.1.45  
langchain-experimental==0.0.57  
langchain-groq==0.1.2  
langchain-openai==0.1.3  
langchain-text-splitters==0.0.1  
langcodes==3.3.0  
langsmith==0.1.49  
chromadb==0.4.24  
ragas==0.1.7  
fastembed==0.2.6
```

Download data

```
! wget "https://arxiv.org/pdf/1810.04805.pdf"
```

Process the PDF Content

```
from langchain.document_loaders import PyPDFLoader
from langchain.text_splitter import RecursiveCharacterTextSplitter
#
loader = PyPDFLoader("1810.04805.pdf")
documents = loader.load()
#
print(len(documents))
```

Perform Native Chunking(RecursiveCharacterTextSplitting)

```
from langchain.text_splitter import RecursiveCharacterTextSplitter

text_splitter = RecursiveCharacterTextSplitter(
    chunk_size=1000,
    chunk_overlap=0,
    length_function=len,
    is_separator_regex=False
)
#
naive_chunks = text_splitter.split_documents(documents)
for chunk in naive_chunks[10:15]:
    print(chunk.page_content+ "\n")

#####RESPONSE#####
BERT BERT
E[CLS] E1 E[SEP] ... ENE1'... EM'
C
T1
T[SEP] ...
TN
T1'...
TM'
[CLS] Tok 1 [SEP] ... Tok NTok 1 ... TokM
Question Paragraph Start/End Span
BERT
E[CLS] E1 E[SEP] ... ENE1'... EM'
C
T1
T[SEP] ...
TN
T1'...
TM'
[CLS] Tok 1 [SEP] ... Tok NTok 1 ... TokM
```

Masked Sentence A Masked Sentence B

Pre-training Fine-Tuning NSP Mask LM Mask LM

Unlabeled Sentence A and B Pair SQuAD

Question Answer Pair NER MNLI Figure 1: Overall pre-training and fine-tuning procedures are used in both pre-training and fine-tuning. The same pre-trained model parameters are used for different downstream tasks. During fine-tuning, all parameters are frozen and a [SEP] token is added in front of every input example, and [SEP] is a special separator token used to separate questions/answers).

ing and auto-encoder objectives have been used for pre-training such models (Howard and Ruder,

2018; Radford et al., 2018; Dai and Le, 2015).

2.3 Transfer Learning from Supervised Data

There has also been work showing effective transfer from supervised tasks with large datasets, such as natural language inference (Conneau et al., 2017) and machine translation (McCann et al., 2017). Computer vision research has also demonstrated the importance of transfer learning from large pre-trained models, where an effective recipe is to fine-tune models pre-trained with ImageNet (Deng et al., 2009; Yosinski et al., 2014).

3 BERT

We introduce BERT and its detailed implementation in this section. There are two steps in our framework: pre-training and fine-tuning. During pre-training, the model is trained on unlabeled data over different pre-training tasks. For fine-tuning, the BERT model is first initialized with the pre-trained parameters, and all of the parameters are fine-tuned using labeled data from the downstream tasks. Each downstream task has sep-

arate fine-tuned models, even though they are initialized with the same pre-trained parameters. The question-answering example in Figure 1 will serve as a running example for this section.

A distinctive feature of BERT is its unified architecture across different tasks. There is minimal difference between the pre-training and the final downstream architecture.

Model Architecture BERT’s model architecture is a multi-layer bidirectional Transformer encoder based on the original implementation described in Vaswani et al. (2017) and released in the `tensorflow/tensorflow` library.¹ Because the use of Transformers has become common and our implementation is almost identical to the original, we will omit an exhaustive background description of the model architecture and refer readers to Vaswani et al. (2017) as well as excellent guides such as “The Annotated Transformer.”²

In this work, we denote the number of layers (i.e., Transformer blocks) as L , the hidden size as

H , and the number of self-attention heads as A .³

We primarily report results on two model sizes:

BERT BASE ($L=12$, $H=768$, $A=12$, Total Parameters=110M) and BERT LARGE ($L=24$, $H=1024$, $A=16$, Total Parameters=340M).

BERT BASE was chosen to have the same model size as OpenAI GPT for comparison purposes.

Critically, however, the BERT Transformer uses bidirectional self-attention, **while** the GPT Transformer uses constrained self-attention where every token can only attend to context to its left.⁴

1<https://github.com/tensorflow/tensor2tensor>
2<http://nlp.seas.harvard.edu/2018/04/03/attention.html>
3In **all** cases we **set** the feed-forward/filter size to be 4H, i.e., **3072** for the H= **768** and **4096** for the H= **1024** .
4We note that **in** the literature the bidirectional Trans-

Input/Output Representations To make BERT handle a variety of down-stream tasks, our **input** representation **is** able to unambiguously represent both a single sentence **and** a pair of sentences (e.g.,(Question, Answer)) **in** one token sequence. Throughout this work, a “sentence” can be an arbitrary span of contiguous text, rather than an actual linguistic sentence. A “sequence” refers to the **in**-put token sequence to BERT, which may be a single sentence **or** two sentences packed together. We use WordPiece embeddings (Wu et al., 2016) **with** a **30,000** token vocabulary. The first token of every sequence **is** always a special classification token ([CLS]). The final hidden state corresponding to this token **is** used **as** the aggregate sequence representation **for** classification tasks. Sentence pairs are packed together into a single sequence. We differentiate the sentences **in** two ways. First, we separate them **with** a special token ([SEP]). Second, we add a learned embed-

Instantiate Embedding Model

```
from langchain_community.embeddings.fastembed import FastEmbedEmbeddings
embed_model = FastEmbedEmbeddings(model_name="BAAI/bge-base-en-v1.5")
```

Setup the API Key for LLM

```
from google.colab import userdata
from groq import Groq
from langchain_groq import ChatGroq
#
groq_api_key = userdata.get("GROQ_API_KEY")
```


Perform Semantic Chunking

We're going to be using the ``percentile`` threshold as an example today — but there's three different strategies you could use on Semantic Chunking):

- ``percentile`` (default) — In this method, all differences between sentences are calculated, and then any difference greater than the X percentile is split.
- ``standard_deviation`` — In this method, any difference greater than X standard deviations is split.
- ``interquartile`` — In this method, the interquartile distance is used to split chunks.

NOTE: This method is currently experimental and is not in a stable final form — expect updates and improvements in the coming months

```
from langchain_experimental.text_splitter import SemanticChunker
from langchain_openai.embeddings import OpenAIEmbeddings

semantic_chunker = SemanticChunker(embed_model, breakpoint_threshold_type="perce
#
semantic_chunks = semantic_chunker.create_documents([d.page_content for d in doc
#
for semantic_chunk in semantic_chunks:
    if "Effect of Pre-training Tasks" in semantic_chunk.page_content:
        print(semantic_chunk.page_content)
        print(len(semantic_chunk.page_content))
```

#####RESPONSE#####

Dev Set

Tasks MNLI-m QNLI MRPC SST-2 SQuAD

(Acc) (Acc) (Acc) (Acc) (F1)

BERT BASE 84.4 88.4 86.7 92.7 88.5

No NSP 83.9 84.9 86.5 92.6 87.9

LTR & No NSP 82.1 84.3 77.5 92.1 77.8

+ BiLSTM 82.1 84.1 75.7 91.6 84.9

Table 5: Ablation over the pre-training tasks using the

BERT BASE architecture. “No NSP” is trained without

the next sentence prediction task. “LTR & No NSP” is

trained as a left-to-right LM without the next sentence

prediction, like OpenAI GPT. “+ BiLSTM” adds a ran-

domly initialized BiLSTM on top of the “LTR + No

NSP” model during fine-tuning. ablation studies can be found in Appendix C. 5.1 E

We demonstrate the importance of the deep bidi-

rectionality of BERT by evaluating two pre-

training objectives using exactly the same pre-

training data, fine-tuning scheme, and hyperparameters as BERT BASE :

No NSP : A bidirectional model which is trained using the “masked LM” (MLM) but without the “next sentence prediction” (NSP) task. LTR & No NSP : A left-context-only model is trained using a standard Left-to-Right (LTR) LM,

Tasks	Dev Set				
	MNLI-m (Acc)	QNLI (Acc)	MRPC (Acc)	SST-2 (Acc)	SQuAD (F1)
BERT _{BASE}	84.4	88.4	86.7	92.7	88.5
No NSP	83.9	84.9	86.5	92.6	87.9
LTR & No NSP	82.1	84.3	77.5	92.1	77.8
+ BiLSTM	82.1	84.1	75.7	91.6	84.9

Table 5: Ablation over the pre-training tasks using the BERT_{BASE} architecture. “No NSP” is trained without the next sentence prediction task. “LTR & No NSP” is trained as a left-to-right LM without the next sentence prediction, like OpenAI GPT. “+ BiLSTM” adds a randomly initialized BiLSTM on top of the “LTR + No NSP” model during fine-tuning.

ablation studies can be found in Appendix C.

5.1 Effect of Pre-training Tasks

We demonstrate the importance of the deep bidirectionality of BERT by evaluating two pre-training objectives using exactly the same pre-training data, fine-tuning scheme, and hyperparameters as BERT_{BASE}:

No NSP: A bidirectional model which is trained using the “masked LM” (MLM) but without the “next sentence prediction” (NSP) task.

LTR & No NSP: A left-context-only model which is trained using a standard Left-to-Right (LTR) LM, rather than an MLM. The left-only constraint was also applied at fine-tuning, because removing it introduced a pre-train/fine-tune mismatch that degraded downstream performance. Additionally, this model was pre-trained without the NSP task. This is directly comparable to OpenAI GPT, but using our larger training dataset, our input representation, and our fine-tuning scheme.

We first examine the impact brought by the NSP task. In Table 5, we show that removing NSP hurts performance significantly on QNLI, MNLI, and SQuAD 1.1. Next, we evaluate the impact of training bidirectional representations by comparing “No NSP” to “LTR & No NSP”. The LTR model performs worse than the MLM model on all tasks, with large drops on MRPC and SQuAD.

For SQuAD it is intuitively clear that a LTR model will perform poorly at token predictions, since the token-level hidden states have no right-side context. In order to make a good faith attempt at strengthening the LTR system, we added a randomly initialized BiLSTM on top. This does significantly improve results on SQuAD, but the

results are still far worse than those of the pre-trained bidirectional models. The BiLSTM hurts performance on the GLUE tasks.

We recognize that it would also be possible to train separate LTR and RTL models and represent each token as the concatenation of the two models, as ELMo does. However: (a) this is twice as expensive as a single bidirectional model; (b) this is non-intuitive for tasks like QA, since the RTL model would not be able to condition the answer on the question; (c) this it is strictly less powerful than a deep bidirectional model, since it can use both left and right context at every layer.

5.2 Effect of Model Size

In this section, we explore the effect of model size on fine-tuning task accuracy. We trained a number of BERT models with a differing number of layers, hidden units, and attention heads, while otherwise using the same hyperparameters and training procedure as described previously.

Results on selected GLUE tasks are shown in Table 6. In this table, we report the average Dev Set accuracy from 5 random restarts of fine-tuning. We can see that larger models lead to a strict accuracy improvement across all four datasets, even for MRPC which only has 3,600 labeled training examples, and is substantially different from the pre-training tasks. It is also perhaps surprising that we are able to achieve such significant improvements on top of models which are already quite large relative to the existing literature. For example, the largest Transformer explored in Vaswani et al. (2017) is (L=6, H=1024, A=16) with 100M parameters for the encoder, and the largest Transformer we have found in the literature is (L=64, H=512, A=2) with 235M parameters (Al-Rfou et al., 2018). By contrast, BERT_{BASE} contains 110M parameters and BERT_{LARGE} contains 340M parameters.

It has long been known that increasing the model size will lead to continual improvements on large-scale tasks such as machine translation and language modeling, which is demonstrated by the LM perplexity of held-out training data shown in Table 6. However, we believe that this is the first work to demonstrate convincingly that scaling to extreme model sizes also leads to large improvements on very small scale tasks, provided that the model has been sufficiently pre-trained. Peters et al. (2018b) presented

Instantiate the Vectorstore

```
hain_community.vectorstores import Chroma
hunk_vectorstore = Chroma.from_documents(semantic_chunks, embedding=embed_model)
```

We will “limit” our semantic retriever to $k = 1$ to demonstrate the power of the semantic chunking strategy while maintaining similar token counts between the semantic and naive retrieved context.

Instantiate Retrieval Step

```
semantic_chunk_retriever = semantic_chunk_vectorstore.as_retriever(search_kwargs={
    "k": 1
})
semantic_chunk_retriever.invoke("Describe the Feature-based Approach with BERT?"

#####RESPONSE#####
[Document(page_content='The right part of the paper represents the\nDev set resu
```

Instantiate Augmentation Step(for content Augmentation)

```
from langchain_core.prompts import ChatPromptTemplate

rag_template = """\
Use the following context to answer the user's query. If you cannot answer, please say so.

User's Query:
{question}

Context:
{context}
"""

rag_prompt = ChatPromptTemplate.from_template(rag_template)
```

Instantiate the Generation Step

```
chat_model = ChatGroq(temperature=0,
                       model_name="mixtral-8x7b-32768",
                       api_key=userdata.get("GROQ_API_KEY"),)
```

Creating a RAG Pipeline Utilizing Semantic Chunking

```

from langchain_core.runnables import RunnablePassthrough
from langchain_core.output_parsers import StrOutputParser

semantic_rag_chain = (
    {"context" : semantic_chunk_retriever, "question" : RunnablePassthrough()}
    | rag_prompt
    | chat_model
    | StrOutputParser()
)

```

Ask Question 1

```

semantic_rag_chain.invoke("Describe the Feature-based Approach with BERT?")

##### RESPONSE #####
The feature-based approach with BERT, as mentioned in the context, involves using BERT.

To use BERT in a feature-based approach, the last 4 layers of BERT are concatenated with the input sequence.

The context also mentions that fine-tuning BERT is surprisingly robust to different input sequences.

In summary, the feature-based approach with BERT involves using the last 4 layers of BERT.

```

Ask Question 2

```

semantic_rag_chain.invoke("What is SQuADv2.0?")

##### RESPONSE #####
SQuAD v2.0, or Squad Two Point Zero, is a version of the Stanford Question Answering Dataset.

```

Ask Question 3

```

semantic_rag_chain.invoke("What is the purpose of Ablation Studies?")

##### RESPONSE #####
Ablation studies are used to understand the impact of different components or sub-components on the overall performance of a model.

```

Implement a RAG pipeline using Naive Chunking Strategy

```
naive_chunk_vectorstore = Chroma.from_documents(naive_chunks, embedding=embed_mo
naive_chunk_retriever = naive_chunk_vectorstore.as_retriever(search_kwargs={"k"
naive_rag_chain = (
    {"context" : naive_chunk_retriever, "question" : RunnablePassthrough()}}
    | rag_prompt
    | chat_model
    | StrOutputParser()
)
```

Note : Here we are going to use $k = 5$;this is to “make it a fair comparison” between the two strategies.

Ask Question 1

```
naive_rag_chain.invoke("Describe the Feature-based Approach with BERT?")

#####RESPONSE#####
The Feature-based Approach with BERT involves extracting fixed features from the
```

Ask Question 2

```
naive_rag_chain.invoke("What is SQuADv2.0?")
#####RESPONSE#####
SQuAD v2.0, or the Stanford Question Answering Dataset version 2.0, is a collect
```

Ask Question 3

```
naive_rag_chain.invoke("What is the purpose of Ablation Studies?")

#####RESPONSE#####
Ablation studies are used to evaluate the effect of different components or sett
```

For instance, one ablation study investigates the effect of the number of training documents on the performance of the model.

Another ablation study focuses on different masking procedures during pre-training.

Ragas Assessment Comparison for Semantic Chunker

split documents using RecursiveCharacterTextSplitter

```
synthetic_data_splitter = RecursiveCharacterTextSplitter(  
    chunk_size=1000,  
    chunk_overlap=0,  
    length_function=len,  
    is_separator_regex=False  
)  
#  
synthetic_data_chunks = synthetic_data_splitter.create_documents([d.page_content  
print(len(synthetic_data_chunks))
```

Create the Following Datasets

- *Questions — synthetically generated (grogq-mixtral-8x7b-32768)*
- *Contexts — created above(Synthetic data chunks)*
- *Ground Truths — synthetically generated (grogq-mixtral-8x7b-32768)*
- *Answers — generated from our Semantic RAG Chain*

```
questions = []  
ground_truths_semantic = []  
contexts = []  
answers = []  
  
question_prompt = """\  
You are a teacher preparing a test. Please create a question that can be answered  
  
Context:  
{context}  
"""  
  
question_prompt = ChatPromptTemplate.from_template(question_prompt)  
  
ground_truth_prompt = """\  
Use the following context and question to answer this question using *only* the
```

```

Question:
{question}

Context:
{context}
"""

ground_truth_prompt = ChatPromptTemplate.from_template(ground_truth_prompt)

question_chain = question_prompt | chat_model | StrOutputParser()
ground_truth_chain = ground_truth_prompt | chat_model | StrOutputParser()

for chunk in synthetic_data_chunks[10:20]:
    questions.append(question_chain.invoke({"context" : chunk.page_content}))
    contexts.append([chunk.page_content])
    ground_truths_semantic.append(ground_truth_chain.invoke({"question" : question
    answers.append(semantic_rag_chain.invoke(questions[-1]))

```

Note: for experimental purpose we have considered only 10 samples

Format the content generated into HuggingFace Dataset Format

```

from datasets import load_dataset, Dataset

qagc_list = []

for question, answer, context, ground_truth in zip(questions, answers, contexts,
    qagc_list.append({
        "question" : question,
        "answer" : answer,
        "contexts" : context,
        "ground_truth" : ground_truth
    })


eval_dataset = Dataset.from_list(qagc_list)
eval_dataset

#####RESPONSE#####
Dataset({
  features: ['question', 'answer', 'contexts', 'ground_truth'],
  num_rows: 10
})

```

Implement Ragas metrics and evaluate our created dataset.

ragas score



```
from ragas.metrics import (
    answer_relevancy,
    faithfulness,
    context_recall,
    context_precision,
)

#
from ragas import evaluate

result = evaluate(
    eval_dataset,
    metrics=[
        context_precision,
        faithfulness,
        answer_relevancy,
        context_recall,
    ],
    llm=chat_model,
    embeddings=embed_model,
    raise_exceptions=False
)
```

Here I had tried to use open source LLM using Groq. But got a rate limit error :

```
limits for more information.', 'type': 'tokens', 'code': 'rate_limit_exceeded'}}
```


groqcloud

Limits

These are the rate limits for your organization

ID	REQUESTS PER MINUTE	REQUESTS PER DAY	TOKENS PER MINUTE
mixtral-8x7b-32768	30	14,400	9,000
gemma-7b-it	30	14,400	15,000
llama3-70b-8192	30	14,400	9,000
llama3-8b-8192	30	14,400	15,000
llama2-70b-4096	30	14,400	15,000

Organization

Billing

Limits

Team

So redirected the LLM to use OpenAI which is by default used in RAGAS framework.

Set up OpenAI API keys

```
import os
from google.colab import userdata
import openai
os.environ['OPENAI_API_KEY'] = userdata.get('OPENAI_API_KEY')
openai.api_key = os.environ['OPENAI_API_KEY']
```

```
from ragas import evaluate

result = evaluate(
    eval_dataset,
    metrics=[
        context_precision,
        faithfulness,
        answer_relevancy,
        context_recall,
    ],
)
result

#####RESPONSE#####
{'context_precision': 1.0000, 'faithfulness': 0.8857, 'answer_relevancy': 0.9172}
```

```
#Extract the details into a dataframe
results_df = result.to_pandas()
results_df
```

	question	answer	contexts	ground_truth	context_precision	faithfulness	answer_relevancy	context_recall
0	Question: What special symbols are added to the input example...	The special symbols added to the input example...	[BERT BERT \n[CLS] E1 E[SEP] ... ENE1' ... EM...	The special symbols added to the input example...	1.0	1.000000	0.910074	1.0
1	Question:\nAccording to the passage, what is o...	One example of a supervised task with a large ...	[2018; Radford et al., 2018; Dai and Le, 2018]...	The passage mentions that transfer learning fr...	1.0	1.000000	0.932776	1.0
2	Question:\nBased on the provided context, how ...	The BERT model architecture has a number of la...	[arate fine-tuned models, even though they are ...	Based on the provided context, the BERT model ...	1.0	1.000000	0.877870	1.0
3	Question:\nWhy did the developers choose BERT ...	The developers chose BERT BASE to have the sam...	[H, and the number of self-attention heads as ...	The developers chose BERT BASE to have the sam...	1.0	1.000000	0.919327	1.0
4	Question: How does the input representation of...	The input representation of BERT unambiguously...	[Input/Output Representations To make BERT\inha...	The input representation of BERT unambiguously...	1.0	1.000000	0.916391	1.0
5	Question: According to the passage, what are t...	The two unsupervised tasks used to pre-train B...	[ding to every token indicating whether it bel...	The two unsupervised tasks used to pre-train B...	1.0	1.000000	0.905318	1.0
6	Question:\nHow can a deep bidirectional repres...	A deep bidirectional representation can be tra...	[standard conditional language models can only...	A deep bidirectional representation can be tra...	1.0	1.000000	0.888184	1.0
7	Question: According to the given context, what...	During training, the masked words are replaced...	[denoising auto-encoders (Vincent et al., 2008...	During training, the masked words are replaced...	1.0	1.000000	0.907590	1.0
8	Question: Based on the given context, what is ...	The purpose of pre-training a model with a bin...	[ence (NLI) are based on understanding the rel...	The purpose of pre-training a model with a bin...	1.0	0.857143	0.944533	1.0
9	Question: Based on the context, what two sourc...	Based on the provided context, the two sources...	[[CLS] helikesplay\ Ing[SEP] mydogiscute[SEP]...	The two sources used for the pre-training corp...	1.0	0.000000	0.966233	1.0

Ragas Assessment Comparison for Naive Chunker

```
import tqdm
questions = []
ground_truths_semantic = []
contexts = []
answers = []
for chunk in tqdm.tqdm(synthetic_data_chunks[10:20]):
    questions.append(question_chain.invoke({"context" : chunk.page_content}))
    contexts.append([chunk.page_content])
    ground_truths_semantic.append(ground_truth_chain.invoke({"question" : question
    answers.append(naive_rag_chain.invoke(questions[-1]))
```

Formulate naive chunking evaluation dataset

```
qagc_list = []

for question, answer, context, ground_truth in zip(questions, answers, contexts,
qagc_list.append({
    "question" : question,
    "answer" : answer,
    "contexts" : context,
    "ground_truth" : ground_truth
}))

naive_eval_dataset = Dataset.from_list(qagc_list)
naive_eval_dataset

#####RESPONSE#####
Dataset({
    features: ['question', 'answer', 'contexts', 'ground_truth'],
```

```
num_rows: 10
})
```

Evaluate our created dataset using RAGAS framework

Open in app ↗



Search

Write



```
sult.to_pandas()
```

```
####RESPONSE #####
```

```
0, 'faithfulness': 0.9500, 'answer_relevancy': 0.9182, 'context_recall': 1.0000}
```

	question	answer	contexts	ground_truth	context_precision	faithfulness	answer_relevancy	context_recall
0	Question: What special symbols are added to the input example...	The special symbols added to the input example...	[BERT BERT [CLS] E1 E[SEP] ... ENET1... EM...	The special symbols added to the input example...	1.0	1.0	0.910074	1.0
1	Question: According to the passage, what is o...	One example of a supervised task with a large ...	[2018; Radford et al., 2018; Dai and Le, 2015]...	The passage mentions that transfer learning fr...	1.0	0.5	0.932371	1.0
2	Question: Based on the provided context, how ...	The BERT model architecture has L layers, wher...	[arate fine-tuned models, even though they are ...	Based on the provided context, the BERT model ...	1.0	1.0	0.878412	1.0
3	Question: Why did the developers choose BERT ...	The developers chose BERT BASE to have the sam...	[H, and the number of self-attention heads as ...	The developers chose BERT BASE to have the sam...	1.0	1.0	0.918327	1.0
4	Question: How does the input representation of...	The input representation of BERT unambiguously...	[Input/Output Representations To make BERT/rha...	The input representation of BERT unambiguously...	1.0	1.0	0.918391	1.0
5	Question: According to the passage, what are L...	The two unsupervised tasks used to pre-train B...	[ding to every token indicating whether it bel...	The two unsupervised tasks used to pre-train B...	1.0	1.0	0.905318	1.0
6	Question: How can a deep bidirectional repres...	A deep bidirectional representation can be tra...	[standard conditional language models can only...	A deep bidirectional representation can be tra...	1.0	1.0	0.888336	1.0
7	Question: According to the given context, what...	During training, the masked words are replaced...	[denoising auto-encoders (Vincent et al., 2006...	During training, the masked words are replaced...	1.0	1.0	0.916917	1.0
8	Question: Based on the given context, what is ...	The purpose of pre-training a model with a bin...	[ence (NLI) are based on understanding the rel...	The purpose of pre-training a model with a bin...	1.0	1.0	0.944533	1.0
9	Question: Based on the context, what two sourc...	The two sources used for the pre-training corp...	[[CLS] helikesplay[[SEP] mydogscute[SEP] ...	The two sources used for the pre-training corp...	1.0	1.0	0.868233	1.0

Conclusion

Here we can see that the results of both Semantic Chunking and Naive Chunking are almost same except that Naive Chunker has a better factual

representation of answer with a score of 0.95 when compared to score of 0.88 of Semantic Chunker.

In conclusion, semantic chunking enables the grouping of contextually similar information, allowing for the creation of independent and meaningful segments. This approach enhances the efficiency and effectiveness of large language models by providing them with focused inputs, ultimately improving their ability to comprehend and process natural language data.

References:

<p>GroqCloud</p> <p>Experience the fastest inference in the world</p> <p>console.groq.com</p>	
<p>Metrics Ragas</p> <p>Skip to content Just like in any machine learning system, the performance of individual components within the LLM and...</p> <p>docs.ragas.io</p>	
<p>Semantic Chunking 🦜🔗 LangChain</p> <p>Splits the text based on semantic similarity.</p> <p>python.langchain.com</p>	



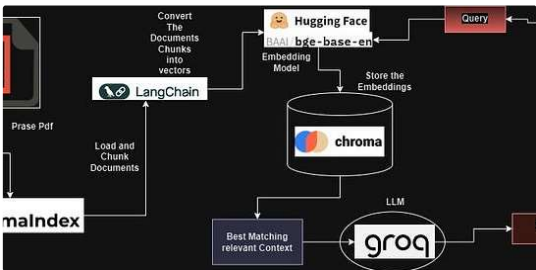
Written by Plaban Nayak

1.93K Followers · Editor for The AI Forum

Machine Learning and Deep Learning enthusiast

Following

More from Plaban Nayak and The AI Forum




Plaban Nayak in The AI Forum

RAG on Complex PDF using LlamaParse, Langchain and Groq

Retrieval-Augmented Generation (RAG) is a new approach that leverages Large Language...

13 min read · Apr 7, 2024

357 3



Plaban Nayak in The AI Forum

Implementing Advanced RAG in Langchain using RAPTOR

Usually in conventional RAG we often rely on retrieving short contiguous text chunks for...

283 min read · Mar 24, 2024

316 1

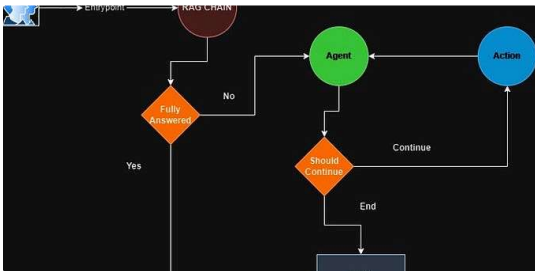
Environment	Cloud Deployment
	<input checked="" type="checkbox"/> (Self-hosted)
	<input checked="" type="checkbox"/> (Managed)
	<input checked="" type="checkbox"/> (Self-hosted)
	<input checked="" type="checkbox"/> (Self-hosted)
	<input checked="" type="checkbox"/> (Self-hosted)
	<input checked="" type="checkbox"/> (Self-hosted)
	<input checked="" type="checkbox"/> (Self-hosted)
	<input checked="" type="checkbox"/> (Self-hosted)

Plaban Nayak in The AI Forum

Which Vector Database Should You Use? Choosing the Best One for...

Introduction

9 min read · Apr 19, 2024



Plaban Nayak in The AI Forum

Implementing Agentic RAG using Langchain

Native RAG

12 min read · Mar 31, 2024

See all from Plaban Nayak

See all from The AI Forum

Recommended from Medium

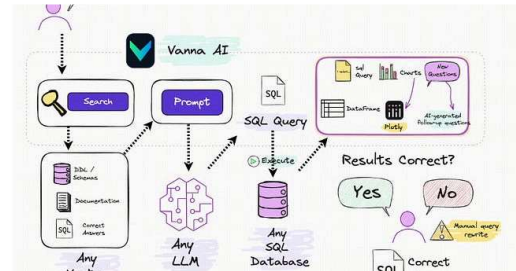


 Aniket Hingane... in Artificial Intelligence in Plain En...

RAG 2.0 : Your AI's Scattered Brain Just Got Organized

Teamwork Makes the Dream Work: The Collaborative Core of RAG 2.0

★ · 9 min read · Apr 23, 2024



 Angelina Yang

No More Text2SQL, It's Now RAG2SQL!

Text2SQL was awfully popular last year. I talked with a vendor looking to buy, and also...

★ · 3 min read · Apr 13, 2024

Lists




Natural Language Processing

1419 stories · 917 saves



Staff Picks

634 stories · 934 saves




Anurag Mishra in CodeX

Build End-to-End RAG Pipeline with Monitoring and Evaluation using...

Introduction

7 min read · Apr 21, 2024

👏 109 💬 📌 ⋮



Vishal Rajput in AIGuys

AI Agents Are All You Need

AI Agents, Understanding the role of Tools, Memory, and Planning in making them work.

★ · 14 min read · Apr 22, 2024

👏 447 💬 1 📌 ⋮



Silvia Onofrei in Towards Data Science

Relation Extraction with Llama3 Models

Enhanced relation extraction by fine-tuning Llama3-8B with a synthetic dataset created...

12 min read · Apr 26, 2024

👏 424 💬 2 📌 ⋮



Fabio Matricardi in The AI Explorer

Microsoft strikes back—Phi-3 may change the entire game.

A Highly Capable Language Model Locally on Your Phone! We have already the GGUF...

★ · 9 min read · Apr 24, 2024

👏 614 💬 3 📌 ⋮

See more recommendations